

AGENTREINS

Passive Monitor → Chief Surgeon

A local-first console stops watching your agents and starts directing them — so the system keeps its conceptual integrity.

4 diagrams


macOS · local-first

Alpha, early and fragile

BEFORE

AgentReins as a Passive Monitor

It watches every agent. It cannot stop, steer, or scope any of them.



Records everything. Controls nothing.

No authority to stop a runaway agent, narrow its scope, or reject an unsafe change.




Every agent. Every socket. Every file — on the record.

Process lineage, file writes and network sockets from every agent, captured to an append-only SQLite WAL.

Many Agents, One Codebase, No Owner

Every agent is competent alone. Together they destroy conceptual integrity.

PARALLEL AGENTS



FAILURE MODES

- Context Collapse**
each agent sees only its own slice
- Hallucinated APIs**
symbols invented, never defined
- Conflicting Edits**
two agents rewrite the same block
- Spaghetti Code**
it compiles, but nobody owns the design

Nobody owns the architecture. Everyone edits it.
This is the exact failure Brooks named in 1975 — and it is back, with agents.




The bottleneck is not model quality. It is ownership.

Context collapse, hallucinated APIs, conflicting edits, spaghetti code.

Brooks' Surgical Team, 1975

From The Mythical Man-Month: one surgeon operates, the team clears the path.




One surgeon owns the design. Everyone else clears the path.

The Surgical Team from The Mythical Man-Month, revived for multi-agent orchestration.

AgentReins as the AI Chief Surgeon

From a monitoring plane to a control plane: observe, decide, delegate, verify, record.



One mind decides. Many hands execute.

Global repo tree, call graph and intent stream become a scoped task envelope your agents execute inside.

Same telemetry. New authority.

ROLE	MONITOR PLANE (TODAY)	CONTROL PLANE (NEXT)
Authority	read-only observer	issues and revokes tasks
Knowledge	per-event telemetry	global repo tree + call graph + intent stream
Unit of work	an event (write, spawn, connect)	a scoped, budgeted task envelope
Failure handling	records it after the fact	blocks it before it lands
Agents' job	each agent decides for itself	each agent executes a narrow slice
Outcome	transparency	transparency + conceptual integrity

—

Transparency first. Then control. Never control without evidence.